

# Statistics 452: Statistical Learning and Prediction

## Chapter 2: Statistical Learning, R supplement

Brad McNeney

# What is R?

- ▶ R is an open-source environment for statistical computing and graphics.
- ▶ Started in the mid-1990's at Auckland University
- ▶ Now maintained by a team of experts called the R Development Core Team
- ▶ A “packages” system allows any user to bundle R code, data and examples together.
  - ▶ Load packages with `library()`
- ▶ R and R packages are distributed through the Comprehensive R Archive Network (CRAN).
- ▶ SFU has a CRAN mirror at <http://cran.stat.sfu.ca>

# What does “environment” mean?

- ▶ R is a fully-functioning programming environment with all the usual constructs, such as
  - ▶ conditionals (if-then-else),
  - ▶ loops
  - ▶ user-defined functions.
- ▶ In addition there are built-in facilities for
  - ▶ data input, storage, manipulation, and output
  - ▶ optimization, matrix computation, etc.,
  - ▶ random number generation,
  - ▶ data analysis and graphics.
- ▶ “Base” R is good, but it is the package system that makes R great.

# Starting R

- ▶ Start R by starting RStudio.
- ▶ The “Console” window is where you can type your commands.
- ▶ However, it is good practice to open an R script, type your commands in the script, and then submit the commands to the R console.
  - ▶ Session -> Set Working Directory to set the working directory
  - ▶ File -> New File -> R Script to open a new R script
  - ▶ type your commands into the script
  - ▶ put your cursor on the line you want to submit and hit Ctrl-enter
- ▶ Save your script for later use.
- ▶ More on the RStudio interface at <https://support.rstudio.com/hc/en-us/sections/200107586-Using-RStudio>

# R Cheatsheets

- ▶ See the RStudio cheatsheets page:  
<https://www.rstudio.com/resources/cheatsheets/>
- ▶ Or use Google to find one that works for you.

# R objects

- ▶ In R, data structures and functions are all referred to as “objects”.
- ▶ Objects are created with the assignment operator `<-`; e.g., `x <- 1`.
  - ▶ The objects a user creates from the R console are contained in the user’s workspace, called the global environment.
  - ▶ Use `ls()` to see a list of all objects in the workspace.
  - ▶ Use `rm(x)` to remove object `x` from the workspace.

# R Data Structures

- ▶ Focus on four common data structures: atomic vectors, lists, matrices and data frames.
- ▶ Atomic vectors and lists are 1d, while matrices and data frames are 2d objects
- ▶ R has no true scalars; e.g., in `x<-1`, `x` is a vector of length one.
- ▶ R also has an array data structure for higher dimensional elements that we will not discuss.
- ▶ Use `str()` to see the structure of an object

# Vectors

- ▶ Vectors can be either atomic or list
  - ▶ atomic vectors must be comprised entirely of logical, integer, double (numeric) or character elements
  - ▶ lists can be comprised of multiple data types
- ▶ Data vectors can be created with `c()` or `list()`:

```
avec <- c(50,200,77)
lvec <- list(50,200,77,c("grey","thin"))
```



# Combining vectors

- Use `c()` to combine vectors

```
c(avec, c(100, 101))
```

```
## [1] 50 200 77 100 101
```

```
c(lvec, TRUE)
```

```
## [[1]]
```

```
## [1] 50
```

```
##
```

```
## [[2]]
```

```
## [1] 200
```

```
##
```

```
## [[3]]
```

```
## [1] 77
```

```
##
```

```
## [[4]]
```

```
## [1] "grey" "thin"
```

```
##
```

```
## [[5]]
```

```
## [1] TRUE
```

# Factors

- ▶ The statistical concept of a factor is important in experimental design.
- ▶ Factors are implemented in R as atomic vectors with attributes class and levels:

```
trt <- factor(c("drug1", "placebo", "placebo", "drug2"))  
attributes(trt)
```

```
## $levels  
## [1] "drug1"    "drug2"    "placebo"  
##  
## $class  
## [1] "factor"
```

```
str(trt)
```

```
## Factor w/ 3 levels "drug1","drug2",...: 1 3 3 2
```

- ▶ The levels are coded numerically (1, 2 and 3) with assigned labels ordered alphabetically ("drug1", "drug2" and "placebo")

## Subsetting vectors and extracting elements

- ▶ Subset with `[` or by name:

```
lvec[c(1,3)] # same as lvec[c("age", "height")]
```

```
## [[1]]
```

```
## [1] 50
```

```
##
```

```
## [[2]]
```

```
## [1] 77
```

- ▶ Extract individual elements with `[[`, or `$` for named objects:

```
lvec[[4]]
```

```
## [1] "grey" "thin"
```

```
lvec$hair
```

```
## NULL
```

# Subsetting and assignment

- ▶ You can combine subsetting and assignment to change the value of vectors

```
avec
```

```
## [1] 50 200 77
```

```
avec[2] <- 210
```

```
avec
```

```
## [1] 50 210 77
```

# Matrices and data frames

- ▶ Though both 2d objects, matrices and data frames are different enough that we will need to discuss them separately.
- ▶ The elements of a matrix must all be of the same type.
- ▶ Data frames are essentially lists where each list element has the same length. Thus data frames can include columns of varying type.

# Matrices

- ▶ Matrices can be created with the `matrix()` function as in

```
A <- matrix(1:4,nrow=2,ncol=2)
A
```

```
##      [,1] [,2]
## [1,]    1    3
## [2,]    2    4
```

- ▶ Here `1:4` is the same as `c(1,2,3,4)`
- ▶ The default is to read the data vector into the matrix column-by-column. To read row-by-row instead use the `byrow=TRUE` argument:

```
A <- matrix(1:4,nrow=2,ncol=2,byrow=TRUE)
A
```

```
##      [,1] [,2]
## [1,]    1    2
## [2,]    3    4
```

# Combining matrices

- Combine matrices with `rbind()` and `cbind()`:

```
rbind(A,matrix(c(5,6),nrow=1,ncol=2))
```

```
##      [,1] [,2]  
## [1,]    1    2  
## [2,]    3    4  
## [3,]    5    6
```

```
cbind(A,A)
```

```
##      [,1] [,2] [,3] [,4]  
## [1,]    1    2    1    2  
## [2,]    3    4    3    4
```

# Subsetting matrices

- ▶ Subset with `[` and a comma to separate rows from columns:

```
A[1,1]
```

```
## [1] 1
```

```
A[1,]
```

```
## [1] 1 2
```

```
A[,1]
```

```
## [1] 1 3
```

- ▶ When a subsetting operation leads to a vector, the dimension of the object is “dropped” from 2 to 1. To prevent this use `drop=FALSE`:

```
A[1,,drop=FALSE]
```

```
##      [,1] [,2]
```

```
## [1,]    1    2
```



## Extracting elements from matrices

```
A[1,1]
```

```
## [1] 1
```

# Data frames

- ▶ Data frames (class `data.frame`) are the usual way to store data in R.
  - ▶ Rows are intended to be observational units, columns variables
  - ▶ Implemented as a list (columns are list elements), but also behave like a matrix in terms of combining and subsetting.
- ▶ Create with `data.frame`:

```
set.seed(1)
n <- 4
x <- 1:n; y <- rnorm(n,mean=x,sd=1) # multiple commands separated by ;
dd <- data.frame(x=x,y=y) # like making a list
str(dd)
```

```
## 'data.frame':    4 obs. of  2 variables:
## $ x: int  1 2 3 4
## $ y: num  0.374 2.184 2.164 5.595
```

# Subsetting and combining data frames

- ▶ Can subset columns like a list:

```
dd$x
```

```
## [1] 1 2 3 4
```

- ▶ Can subset columns/rows and combine like matrices; e.g.,

```
dd[1:2,]
```

```
##      x      y
## 1 1 0.3735462
## 2 2 2.1836433
```

```
zz = data.frame(z=runif(4))
cbind(dd,zz)
```

```
##      x      y      z
## 1 1 0.3735462 0.62911404
## 2 2 2.1836433 0.06178627
## 3 3 2.1643714 0.20597457
## 4 4 5.5952808 0.17655675
```

# Logical operators

- ▶ ! is NOT
- ▶ & and && are AND, with & acting vector-wise and && acting on scalars
- ▶ | and || are OR, with | acting vector-wise and || acting on scalars
- ▶ Make sure you understand the following:

```
x <- c(TRUE,TRUE,FALSE); y <- c(FALSE,TRUE,TRUE)
!x ; x&y ; x&&y ; x|y ; x||y
```

```
## [1] FALSE FALSE TRUE
```

```
## [1] FALSE TRUE FALSE
```

```
## [1] FALSE
```

```
## [1] TRUE TRUE TRUE
```

```
## [1] TRUE
```

# Relational operators

- ▶ Relational operators can be used to compare values in atomic vectors
  - ▶ See `help("Comparison")`
- ▶ `>` is greater than, `>=` is greater than or equal
- ▶ `<` is less than, `<=` is less than or equal
- ▶ `==` is equal and `!=` is not equal
- ▶ Make sure you understand the following:

```
x <- 1:3; y <- 3:1
x>y ; x>=y ; x<y ; x<=y ; x==y ; x!=y
```

```
## [1] FALSE FALSE  TRUE
```

```
## [1] FALSE  TRUE  TRUE
```

```
## [1]  TRUE FALSE FALSE
```

```
## [1]  TRUE  TRUE FALSE
```

```
## [1] FALSE  TRUE FALSE
```

```
## [1]  TRUE FALSE  TRUE
```

# Subsetting vectors with logical expressions

- Can subset with logicals and [:

```
avec
```

```
## [1] 50 210 77
```

```
avec > 100
```

```
## [1] FALSE TRUE FALSE
```

```
avec[avec > 100]
```

```
## [1] 210
```

```
avec[avec > 50 & avec < 100]
```

```
## [1] 77
```

## Subsetting matrices with logical expressions

- Can also subset matrices, but results may not be as expected:

```
A
```

```
##      [,1] [,2]  
## [1,]    1    2  
## [2,]    3    4
```

```
A>1
```

```
##      [,1] [,2]  
## [1,] FALSE TRUE  
## [2,]  TRUE TRUE
```

```
A[A>1] # coerces to a vector
```

```
## [1] 3 2 4
```

# Missing values

- ▶ R has a special data code for missing data: NA
- ▶ Test for and set missing values with `is.na()`

```
avec
```

```
## [1] 50 210 77
```

```
is.na(avec)
```

```
## [1] FALSE FALSE FALSE
```

```
is.na(avec) <- 2
```

```
avec
```

```
## [1] 50 NA 77
```



## R functions: Example

```
f <- function(x) {  
  return(x^2)  
}  
f
```

```
## function(x) {  
##   return(x^2)  
## }
```

## Reading Data: Native format

- ▶ Use `save()` to save R objects to an “R Data” file.
  - ▶ `save.image()` is short-hand to save all objects in the workspace

```
x <- rnorm(100); y <- list(a=1,x=x)
save(x,y,file="test.RData") # Or .rda, or ...
```

- ▶ Load R Data files into the workspace with `load()`.

```
load("test.RData")
file.remove("test.RData")
```

```
## [1] TRUE
```

# Reading Table Format Files

- ▶ `read.table()` is the main function for reading tabular data from plain-text files.
  - ▶ `read.csv()` and `read.delim()` are basically `read.table()` with defaults for reading comma- and tab- delimited files.
- ▶ `write.table()`, `write.csv()` and `write.delim()` are the analogous functions for writing tabular data

```
write.table(matrix(1:9,3,3),file="test.txt")
test <- read.table("test.txt")
file.remove("test.txt")
```

```
## [1] TRUE
```

```
test
```

```
##   V1 V2 V3
## 1  1  4  7
## 2  2  5  8
## 3  3  6  9
```

# Reading files from a URL

- `load()`, `read.table()`, etc. can read data from a URL.

```
baseURL <- "http://people.stat.sfu.ca/~mcneney/Teaching/Stat452/"
rdURL <- url(paste0(baseURL, "Data/PorschePrice.rda"))
load(rdURL)
head(PorschePrice)
```

```
##   Price Age Mileage
## 1  69.4   3   21.5
## 2  56.9   3   43.0
## 3  49.9   2   19.9
## 4  47.4   4   36.0
## 5  42.9   4   44.0
## 6  36.9   6   49.8
```

```
csvURL <- url(paste0(baseURL, "Data/PorschePrice.csv"))
PorschePrice <- read.csv(csvURL)
```

## stringsAsFactors

- ▶ Reading columns that include characters in as factors is controlled by a global option in your R session called `stringsAsFactors`, set to `TRUE` by default.
- ▶ If you want to set to `FALSE` for an R session type `options(stringsAsFactors = FALSE)` into the Console.
- ▶ An alternative is to over-ride the default in the call to `read.table()`:

```
exURL <- url(paste0(baseUrl,"Data/Ex1_1_4.txt"))
ex2 <- read.table(exURL,header=TRUE,sep="\t",
                  stringsAsFactors=FALSE)
```

## Viewing Data: `print()`, `View()` and `edit()`

- ▶ `print()` prints R objects
  - ▶ This function is “generic”, meaning that it will try to find the specific function to print specific objects (e.g., `print.data.frame`).
- ▶ `View()` launches a new window (or RStudio tab) to view a data frame and `edit()` launches a data editor.

# Graphics

- ▶ “Base” graphics in R is good, but `ggplot()` is better.
- ▶ We could spend a lot of time on `ggplot()`, but will just learn what we need as we go.
- ▶ ggplot2 cheatsheet at [<https://www.rstudio.com/wp-content/uploads/2016/11/ggplot2-cheatsheet-2.1.pdf>]
- ▶ Wickham (2009) ggplot2: Elegant graphics for data analysis, Chapters 4 and 5.
- ▶ Chang (2012) R graphics cookbook. Available at [<http://www.cookbook-r.com/Graphs/>]