

Statistics 452: Statistical Learning and Prediction

Chapter 10, part 1: Introduction to Unsupervised Learning

Brad McNeney

2018-11-14

Supervised *versus* Unsupervised Learning


- ▶ Supervised means that there is an outcome y , unsupervised means there is not.
- ▶ Supervised learning has well-defined goals like prediction.
 - ▶ Can check the fitted model by seeing how well it predicts test observations.
- ▶ Unsupervised learning is more exploratory, without an obvious goal.
 - ▶ A common theme is trying to identify simple structure underlying the feature data.
 - ▶ We will discuss dimension reduction by principal components analysis (PCA) and clustering.

Principal Components Analysis (PCA)

- ▶ Goal is low-rank approximation of the X data matrix
 - ▶ Discussed in Chapter 6 and reviewed below.
- ▶ Think of principal components (PCs) as new coordinates for the data vectors.
 - ▶ The first PC is the direction of greatest variation,
 - ▶ The second PC is the direction of second-greatest variation, orthogonal to the first,
 - ▶ And so on.

PCs for Advertising Data

- Text Figure 6.14: The green line is the first PC, the blue line the second.



PCs as Linear Combinations of X 's

- ▶ The details of how the linear combinations are derived are discussed in the text.
- ▶ In the advertising example, the first PC is



$$Z_1 = 0.838X_1 + 0.544X_2$$

where X_1 is population centred by its mean and X_2 is advertising expenditure centred by its mean.

- ▶ The coefficients of the linear combination, $\phi_{11} = 0.838$ and $\phi_{12} = 0.544$, are called the first principal component *loadings*.


Principal Component Scores

- ▶ Projecting each point onto the PCs gives the PC scores.
 - ▶ Projecting a data vector onto a line means finding the point on the line closest to the vector.
- ▶ Text Figure 6.15: Black x's are the first PC score for each observation, distance of each purple dot from the green line is the second PC score.



High-Dimensional Example: Genes Reflect Geography

- First 2 PCs from 197,146 genetic markers on 1,387 European individuals (Novembre *et al.* 2008)



US Arrests Data

- ▶ Dataset that comes with R.
- ▶ From the help file: “This data set contains statistics, in arrests per 100,000 residents for assault, murder, and rape in each of the 50 US states in 1973. Also given is the percent of the population living in urban areas.”

```
data(USArrests) # help(USArrests)
head(USArrests)
```

##		Murder	Assault	UrbanPop	Rape
##	Alabama	13.2	236	58	21.2
##	Alaska	10.0	263	48	44.5
##	Arizona	8.1	294	80	31.0
##	Arkansas	8.8	190	50	19.5
##	California	9.0	276	91	40.6
##	Colorado	7.9	204	78	38.7


```
pcout <- prcomp(USArrests,scale=TRUE)
pcout$rotation # loadings
```

##		PC1	PC2	PC3	PC4
##	Murder	-0.5358995	0.4181809	-0.3412327	0.64922780
##	Assault	-0.5831836	0.1879856	-0.2681484	-0.74340748
##	UrbanPop	-0.2781909	-0.8728062	-0.3780158	0.13387773
##	Rape	-0.5434321	-0.1673186	0.8177779	0.08902432


```
head(pcout$x) # scores
```

##		PC1	PC2	PC3	PC4
##	Alabama	-0.9756604	1.1220012	-0.43980366	0.154696581
##	Alaska	-1.9305379	1.0624269	2.01950027	-0.434175454
##	Arizona	-1.7454429	-0.7384595	0.05423025	-0.826264240
##	Arkansas	0.1399989	1.1085423	0.11342217	-0.180973554
##	California	-2.4986128	-1.5274267	0.59254100	-0.338559240
##	Colorado	-1.4993407	-0.9776297	1.08400162	0.001450164

Scree Plot

- ▶ A scree plot shows the variance (or proportion of total variance) in the direction of each PC.
- ▶ If the variance drops and then levels out, the “elbow” where it levels out is a reasonable choice for a reduced number of PCs that captures most of the variation in the **X**.

```
screeplot(pcout) # or just plot(pcout)
```



- ▶ No obvious elbow.

Iris Data


```
data(iris) # help(iris)
head(iris)
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
## 1	5.1	3.5	1.4	0.2	setosa
## 2	4.9	3.0	1.4	0.2	setosa
## 3	4.7	3.2	1.3	0.2	setosa
## 4	4.6	3.1	1.5	0.2	setosa
## 5	5.0	3.6	1.4	0.2	setosa
## 6	5.4	3.9	1.7	0.4	setosa

```
pcout.i <- prcomp(iris[, -5], scale=TRUE)
```

- ▶ For the iris data, two PCs appear to explain most of the variation.

```
screeplot(pcout.i)
```



Interpretation of Loadings

- ▶ The first PC is a contrast between sepal width and the other variables.
- ▶ The second PC is a weighted average of sepal length and width.


```
pcout.i$rotation
```

##		PC1	PC2	PC3	PC4
##	Sepal.Length	0.5210659	-0.37741762	0.7195664	0.2612863
##	Sepal.Width	-0.2693474	-0.92329566	-0.2443818	-0.1235096
##	Petal.Length	0.5804131	-0.02449161	-0.1421264	-0.8014492
##	Petal.Width	0.5648565	-0.06694199	-0.6342727	0.5235971

Biplot of First Two PCs


- ▶ We can visualize the first two PCs on a scatterplot.
- ▶ A biplot shows
 1. the PC scores for observational units (see left and bottom axes), and
 2. the loadings of the features that define the first two PCs (see top and right axes)

```
biplot(pcout.i,cex=.5,scale=0) #scale=0 avoids scaling of points on plot
```



Biplot of US Arrests Data

```
biplot(pcout, cex=.5, scale=0) #scale=0 avoids scaling of points on plot
```



- ▶ Note different appearance from text: PCs are only unique up to a sign change