

# Statistics 452: Statistical Learning and Prediction

## Chapter 7, Part 3: Smoothing Splines, Local Regression

Brad McNeney

2018-10-23

# Smoothing Splines

# Smoothing Splines Overview

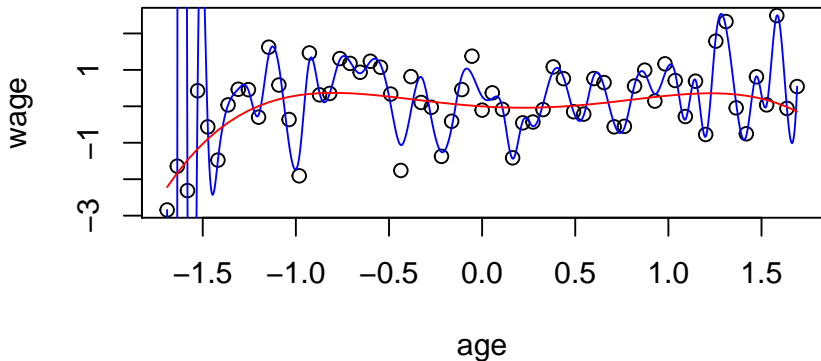
- ▶ Smoothing splines are an alternative approach to devising a smooth curve.
- ▶ Though the derivation is different, we end up with a natural cubic spline, with knots *at each observed data point*.
  - ▶ The estimated coefficients of the basis functions using least squares would interpolate the data points.
  - ▶ To avoid over-fitting, a shrinkage penalty is added to the RSS. The penalty is comprised of a tuning parameter times a penalty term.

# Smoothness

- ▶ We seek a *smooth* function  $g(x)$  that fits the data well.
  - ▶ Want  $RSS = \sum_{i=1}^n (y_i - g(x_i))^2$  small.
- ▶ By smooth we mean not too “rough” (wiggly)
  - ▶ If RSS is the objective function, the curve  $g$  will interpolate the data and will be very rough.
  - ▶ Illustrate by fitting a natural cubic spline to simulated data on “age” and “wage”

```
set.seed(1)
age <- scale(18:80); betas <- c(2,-2,2,-2)
f <- poly(age,degree=4)%*% betas
wage <- f + rnorm(length(age))
sfrit0 <- smooth.spline(age,wage,lambda=0)
newage<-seq(from=min(age),to=max(age),length=1000)
pwage0 <- predict(sfrit0,newage)
```

```
plot(age,wage)
lines(pwage0$x,pwage0$y,col="blue") # fitted curve
lines(age,f,col="red") # true curve
```




# Penalized Objective Function

- ▶ The criterion function to minimize is

$$\sum_{i=1} (y_i - g(x_i))^2 + \lambda \int g''(t)^2 dt. \quad (1)$$

- ▶ Like ridge regression or the lasso, the objective function is of the form  $\text{RSS} + \text{penalty}$ , and we have a tuning, or smoothing parameter  $\lambda$ , but the penalty function is now  $\int g''(t)dt$ .
  - ▶  $g''(t)$  is the curvature of  $g$  at  $t$ .
  - ▶  $\int g''(t)^2 dt$  is a measure of the total curvature.
- ▶ It can be shown that the minimizer of (1) is a natural cubic spline with knots at the unique observed data points, with the coefficients of the basis functions *shrunk* towards zero.
  - ▶ The degree of shrinkage is controlled by  $\lambda$ .

```
sfit1 <- smooth.spline(age,wage,spar=1) #spar= c1+c2*log(lambda)
pwage1 <- predict(sfit1,newage)
plot(age,wage)
lines(pwage0$x,pwage0$y,col="blue") # fitted curve with lambda=0
lines(pwage1$x,pwage1$y,col="green")
lines(age,f,col="red") # true curve
```



# Choosing the Smoothing Parameter

- ▶ Could use cross validation to select  $\lambda$ .
  - ▶ How many folds?
- ▶ It turns out there is a very simple formula for the CV estimate of test error when using leave-out-one CV (LOOCV):

$$CV_{(n)} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{g}_{\lambda}(x_i)^{(-i)})^2 = \frac{1}{n} \sum_{i=1}^n \left( \frac{y_i - \hat{g}_{\lambda}(x_i)}{1 - \{S_{\lambda}\}_{ii}} \right)^2 \quad (2)$$

where

- ▶  $\hat{g}_{\lambda}(x_i)^{(-i)}$  is the fitted value for  $x_i$  when  $(x_i, y_i)$  is held out of the training set,
- ▶  $\hat{g}_{\lambda}(x_i)$  is the fitted value for  $x_i$  when all of the data are used to fit  $g$ ,
- ▶  $\{S_{\lambda}\}_{ii}$  is the  $i$ th diagonal entry of the “smoother matrix”.



# The Smoother Matrix

- ▶ The smoother matrix  $S_\lambda$  is the matrix that turns the response  $y$  into the smooth  $\hat{g}_\lambda$ ; i.e.,  $\hat{g}_\lambda(x) = S_\lambda y$ , where  $x$  is the vector of explanatory variable values and  $\hat{g}_\lambda(x)$  is the vector of  $\hat{g}_\lambda(x_i)$  values.
- ▶ In linear regression, the smoother matrix is called the hat matrix and is denoted  $H$ .
  - ▶  $H$  turns the response  $y$  into  $\hat{y}$ ; i.e.,  $\hat{y} = Hy$ ,
  - ▶ The sum of the diagonal elements of  $H$  is the number of regression coefficients  $p + 1$ .
  - ▶ The diagonal elements are the hat values, or leverages  $h_i$ .
- ▶ Thus the  $CV_{(n)}$  estimate for the smoothing spline is analogous to least squares, with the diagonal elements of  $H$  replaced by those of  $S_\lambda$ .


# Effective Degrees of Freedom

- ▶ In linear regression, the model df  $p + 1$  can be shown to equal the sum of the diagonal elements of  $H$ .
- ▶ By analogy, the sum of the diagonal elements of the smoother matrix is referred to as the *effective degrees of freedom* and is denoted  $df_\lambda$ .
  - ▶ One can show that as  $\lambda$  ranges from 0 to  $\infty$ ,  $df_\lambda$  ranges from  $n$  to 2.
  - ▶ Can treat  $df_\lambda$  as the smoothing parameter and select its value by CV.

```
sfitCV <- smooth.spline(age,wage,cv=TRUE)
sfitCV$df
```

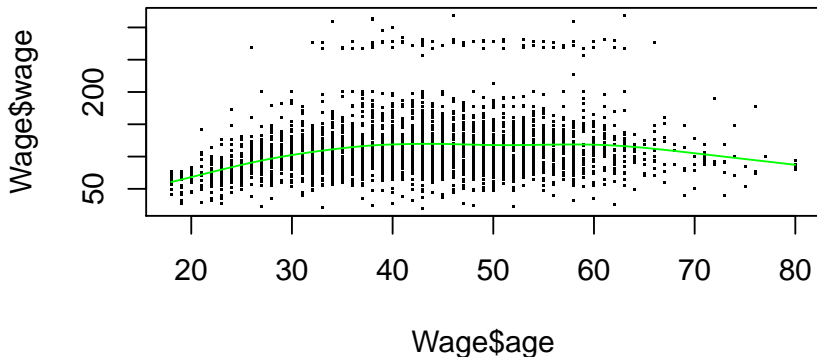
```
## [1] 5.769452
```

```
pwageCV <- predict(sfitCV,newage)
plot(age,wage)
lines(pwageCV$x,pwageCV$y,col="green")
lines(age,f,col="red") # true curve
```



## Example

```
library(ISLR); data(Wage)
plot(Wage$age,Wage$wage,pch=".")
sfit <- smooth.spline(Wage$age,Wage$wage,cv=TRUE)
pwage <- predict(sfit)
lines(sfit,col="green") # can plot output of smooth.spline directly
```



## Local Regression

# Local Regression

- ▶ A variation on KNN: instead of fitting a constant over a neighborhood, fit a weighted regression.
  - ▶ weighted means points in the neighborhood closest to the point of prediction are up-weighted.
  - ▶ the regression could be constant, linear or quadratic.
- ▶ The neighborhood size is referred to as the “span”  $s$ .

# Local Linear Regression Algorithm

- ▶ Select a span  $s$  and a weight function  $K(x_i, x_0)$ .
- ▶ For  $X = x_0$ :
  1. Extract the nearest  $s * n$  points to  $x_0$  to train the local regression.
  2. Assign weight  $K_{0i} = K(x_i, x_0)$  to each neighbor.
  3. Fit a weighted least-squares regression; that is, find the  $\hat{\beta}_0$  and  $\hat{\beta}_1$  that minimize

$$\sum_{i=1}^n K(x_i, x_0)(y_i - [\beta_0 + \beta_1 x_i])^2$$

4.  $\hat{f}(x_0) = \hat{\beta}_0 + \hat{\beta}_1 x_0$ .


```

plotfit <- function(fit,dat,newdat){
  preds <- data.frame(newdat,
    predict(fit,newdata=newdat,se=TRUE))
  preds <- mutate(preds, # approx. CIs from the SEs
    lwr = preds$fit-2*preds$se,upr = preds$fit+2*preds$se)
  ggplot(dat,aes(x=age,y=wage)) + geom_point(alpha=0.1) +
    geom_ribbon(aes(x=age,y=fit,ymin=lwr,ymax=upr),
      data=preds,fill="blue",alpha=.2) +
    geom_line(aes(y=fit),data=preds,color="blue")
}


```



```
library(ggplot2); library(dplyr)
newdat <- data.frame(age=seq(from=min(Wage$age),to=max(Wage$age),length=100))
fit2 <- loess(wage~age,span=0.2,data=Wage)
plotfit(fit2,Wage,newdat)
```



```
fit5 <- loess(wage~age,span=0.5,data=Wage)  
plotfit(fit5,Wage,newdat)
```



# Extensions to Higher Dimension

- ▶ For  $p > 1$  explanatory variables, we can generalize local regression.
  - ▶ Choose neighborhoods based on  $X_1, \dots, X_p$ .
  - ▶ Fit a multiple linear regression.
- ▶ However, it has been observed that local regression performs poorly for  $p > 3$  because there are few points close to  $x_0$  (recall homework 2).